

BARCELONA COMPUTATIONAL FOUNDATION

Working Paper WP0195

Pillar: Kolmogorov Theory (KT)

A Lean 4 Formalization of Kolmogorov Theory

Axioms, Theorems, and Project Status

Machine-checking the KT corollaries against an explicit AIT interface

Giulio Ruffini with Claude Opus 4.8

Barcelona Computational Foundation · Neuroelectrics, Barcelona

`giulio.ruffini@bcom.one`

July 2026

Abstract

We are building **KTAIT**, a Lean 4 + Mathlib formalization whose purpose is narrow and load-bearing: to *machine-check that the Kolmogorov-Theory (KT) corollaries follow logically from an explicit algorithmic-information-theory (AIT) interface*, and that the KT ontology is typed so that structural errors (e.g. comparing a whole pattern to its own part) cannot silently compile. We do *not* re-prove classical AIT: the standard AIT and probability results (the invariance theorem, the coding theorem, symmetry of information, Kraft’s inequality, Bayes’ rule, the Solomonoff–Levin universal semimeasure, the recursion theorem, Rice’s and Chaitin’s theorems, and the Vereshchagin–Vitányi structure-function results) are taken as a clearly delimited *axiom layer*. Every *KT* corollary is then *proved* from that layer with no **sorry**, and a toy model witnesses that the axiom layer is jointly satisfiable (so the corollaries are not vacuously true). This working paper states the methodology, fixes the axiom layer, inventories every theorem/lemma across the KT corpus with its current formalization status, and sets a prioritized roadmap.

Keywords: Kolmogorov Theory, algorithmic information theory, Lean 4, Mathlib, formal verification, algorithmic regulator theorem, persistence, self-model.

Classification: MSC 68Q30, 03B35, 68V15 / ACM F.4.1

Contents

1	Introduction and methodology	4
1.1	The one rule	4
1.2	Why this is worth doing	4
1.3	Soundness: named hypotheses, not global axioms	4
1.4	Scope	4
2	The standard AIT / probability axiom layer	4
3	Current status of the Lean development	5
4	Theorem inventory and formalization status	8
5	Roadmap	10
6	Conclusion	11

1 Introduction and methodology

KT has matured to the point where the dominant risk is no longer in the classical results it invokes, but in whether the *new* KT corollaries are *correctly typed* and *logically entailed* by those results. This project addresses exactly that risk with a proof assistant.¹

1.1 The one rule

The one rule

Standard AIT (and probability) theorems are allowed to be *axioms*. KT corollaries must be *proved* from those axioms with no **sorry**.

If a KT corollary needs a **sorry**, that is a real gap in the paper. If a classical AIT fact is an axiom, that is honest and expected: we are checking the *dependency structure*, not the AIT foundations.

1.2 Why this is worth doing

The bug class we most want to exclude is *ontological*. For example, writing $I_K(A : \mathcal{S})$ where $\mathcal{S} \subset A$ is structurally vacuous (a whole versus its own part); the correct object is *temporal*, $I_K(\mathcal{S}_t : \mathcal{S}_{t+\tau})$ or $K(\mathcal{S}_{t+\tau} \mid \mathcal{S}_t, E, C) \ll K(\mathcal{S}_{t+\tau} \mid C)$. In KTAIT the KT roles (*substrate*, *pattern*, *self-code*, *readout*, *regulator*, *time*) are *distinct types*; the only bridge to a common carrier is an explicit projection, so the bad version *fails to type-check*. The whole-vs-part error becomes a typed obligation, not a silent mistake.

1.3 Soundness: named hypotheses, not global axioms

A naïve encoding axiom $F : \text{forall } (F : \text{AITFrame}), P F$ asserting a law for *every* frame is *inconsistent* here: one can hand-build a frame violating it (e.g. `slack=0` with mismatched K) and derive `False`, which would make all corollaries vacuously true. We therefore state each AIT fact as a *named hypothesis* (a `Prop` about a frame) and let each corollary consume it. Consequently `#print axioms` on every KT corollary shows only Lean’s core axioms (`propext`, `Classical.choice`, `Quot.sound`) — strictly stronger than “rests on named AIT axioms.” The toy model (§3) then exhibits a concrete frame satisfying the hypotheses, witnessing joint satisfiability.

1.4 Scope

The current effort is **Level 1** (a typed, axiomatic KT/AIT interface) plus **Level 1.5** (the faithful *probabilistic* ART, including its Bayesian posterior and the coding-theorem bridge). Out of scope for now: a computability-backed K from a concrete universal machine (Level 2), and the differential-geometric (Lie-group / Noether) layer of the world-models paper,⁴ which is a separate track.

2 The standard AIT / probability axiom layer

These are the results we *assume*. They are classical and not in doubt; collecting them here fixes precisely what the KT proofs are allowed to lean on. (Per K’s directive, starting from these as axioms — and similarly for Bayes — is acceptable and intended.)

Axiom layer A — complexity

Invariance theorem. K is machine-independent up to an additive $O(1)$.

Coding theorem (unconditional). $c_1 2^{-K(x)} \leq m(x) \leq c_2 2^{-K(x)}$, i.e. $-\log_2 m(x) = K(x) \pm O(1)$.

Coding theorem (conditional). $m(x | y) \asymp 2^{-K(x|y)}$.

Kraft–McMillan. $\sum_p 2^{-|p|} \leq 1$, hence $\sum_x m(x) \leq 1$.

Symmetry of information. $I_K(x:y) = K(x) + K(y) - K(x, y) = K(x) - K(x | y^*) + O(\log)$; chain rule $K(x, y) = K(x) + K(y | x^*) + O(1)$.

Multiplicity compression. $N_{\leq L}(x) \geq 2^r \Rightarrow K(x) \leq L - r + O(\log L)$.

Axiom layer B — probability / Bayes

Solomonoff–Levin universal semimeasure. $m(x) = \sum_{U(p)=x} 2^{-|p|}$; conditional $m(R | W) \asymp 2^{-K(R|W)}$.

Bayes' rule + deterministic likelihood. prior $P(p) = 2^{-|p|}$, likelihood $P(x | p) = 1\{U(p) = x\}$, evidence $P(x) = m(x)$.

Excess-length decay. $\Pr\{|p| \geq K(x) + k | x\} \leq 2C 2^{-k}$.

Axiom layer C — computability / logic

Kleene's second recursion theorem (quines); **Rice's theorem**; **Chaitin's incompleteness** (T proves no $K(x) > n$ for $n > c_T$; K uncomputable); **Vereshchagin–Vitányi**: the Kolmogorov structure function $h_x(\alpha)$ is uncomputable and minimal sufficient statistics are not locatable.

3 Current status of the Lean development

The repository (<https://github.com/giulioruffini/KTAIT>) builds green on Lean 4.31.0 + Mathlib v4.31.0. Modules:

Module	Content	Key results
Ontology.lean	six typed KT roles + part-whole guard	SelfCode carries isSub
Basic.lean	AITFrame, IK, NMAI, condStar	symmetry as named Prop
Probability.lean	prefix machine, Bayes posterior, coding thm	Lemma 1 (bridge)
ART.lean	AITProb (posterior), Theorem 2	probabilistic regulator thm
Persistence.lean	Pers, TemporalSelfInfo	persistence lemmas
SelfModel.lean	DeltaSelf, Proposition 3	temporal self-model
SelfModelLimits.lean	quine floor, diagonalization, Chaitin ceiling	WP0192 Principle 1
RegulatorSelection.lean	ChainRule	WP0162 Prop. 1 (regulator selection)
CoarseGraining.lean	RegSuff, structure function	WP0193 Thms A/B; WP0007 emergence
Contrast.lean	contrast fibre posterior	WP0162 Q3 (Eq. 26)
OrbitLabel.lean	conserved orbit labels	WP0162 App. C
NoetherFlow.lean	trajLabel, flow symmetries	generalized Noether (abstract core)
Decoder.lean	Achievable, Inverts (inverse compiler)	WP0058 Prop. 1 (axiom-free)
WriteBack.lean	bandwidth, Darwinian, Lamarckian	WP0058 Prop. 1 (AIT half), Prop. 2
ToyModel.lean	Toy, ToyMachine, ToyGap	satisfiability witnesses
BadStatements.lean	guards bite	part-whole + y vs y^*

Proved KT corollaries (no sorry; #print axioms = Lean core only):

- AITProb.probabilistic_regulator_theorem — ART Theorem 2, $\Pr((W, R) | x, E) \leq C 2^{M(W:R)} 2^{-\Delta}$.

- `AITProb.wrapper_bound` — ART Eq. (6), now *derived* from the coding theorem: $\text{post}(e, x) \leq (1/c_1)2^{K(x)-K(e)}$.
- `AITProb.probabilistic_regulator_theorem_sharp` — ART Theorem 2 *sharp form* retaining the regulator cost: $\text{post} \leq C 2^{M(W:R)} 2^{-\Delta} 2^{-K(R)}$ (no extra hypothesis; the headline drops $2^{-K(R)} \leq 1$ via `probabilistic_regulator_theorem_of_sharp`).
- `AITProb.theorem1_posterior_tilt` — ART Theorem 1, two-sided posterior tilt.
- `AITProb.theorem3_onoff_evidence` — ART Theorem 3, ON/OFF evidence ratio $\asymp 2^\Delta$ (log-free form).
- `CoarseGraining.theoremB / corollaryB` — WP0193 Theorem B / Corollary B: regulatory coarse-graining is uncomputable, by reduction to Vereshchagin–Vitányi.
- `quine_floor`, `self_prediction_dichotomy`, `chaitin_blocks_minimality` — WP0192 Principle 1 / WP0162 Proposition 4: the three independent obstructions to a complete, exact, certifiable self-model (the middle one is an axiom-free diagonalization).
- `contrast_posterior_ranks_by_complexity` — WP0162 Q3 (Eq. 26): the contrast-fiber posterior ranks by joint simplicity, with no $2^{-\Delta}$ tilt.
- `OrbitLabel.genEnergy_conserved` — WP0162 App. C: generalized energy as a conserved orbit label (axiom-free).
- `boundary_sufficient` — WP0162 App. E: the thin boundary as a compressed sufficient statistic.
- `low_complexity_shrinkage` — ART Theorem A4; `CoarseGraining.theoremA`, `CoarseGraining.existence` — WP0193 Theorem A and Proposition 1.
- `Noether.trajLabel_conserved`, `Noether.conserved_comp_symm` — the abstract core of generalized Noether (Structured Dynamics paper): the trajectory label is conserved by the flow, and symmetries commuting with the flow act on conserved quantities. (Continuous-time analogue of the orbit-label theorem.)
- `Homogeneous.homogeneousSpace` — a transitive symmetry action makes the state space a homogeneous space, $X \simeq G/H_x$ (Theorem A4’s algebraic core; the manifold/dimension layer remains future work).
- `Decoder.achievable_computable_of_inverse / Decoder.no_universal_decoder` — WP0058 Proposition 1: a *total* computable inverse compiler that recognizes its own domain would decide membership in the achievable set of the developmental map; hence none exists, and every realized Lamarckian write-back channel is a partial map whose validity domain is fixed in advance. Notable as the one KT corollary here that assumes *no* AIT fact at all — the engine (`achievable_computable_of_inverse`) is unconditional, and only the non-vacuity of the undecidability hypothesis enters as a named Prop (`ExistsUndecidableAchievableSet`). Satisfiability witness: `inverts_id`.
- `WriteBack.bandwidth_le_cond`, the data-processing lemma, and the regime corollaries — WP0058 Proposition 2. With $\lambda_B := I_K(a : H' | H)$ the write-back bandwidth, the master bound $\lambda_B \leq K(H' | H) + O(\log)$ holds but is weak in the Darwinian regime, where undirected mutation makes $K(H' | H)$ large without transmitting anything acquired. The sharp statement is data processing (`condIK_le_condRight`): since a reaches H' only through a channel variable z , $\lambda_B \leq I_K(a : z | H) \leq K(z | H) + O(\log)$. Crucially, the write-back

bandwidth is *not* $I_K(a : H' | H)$: that quantity is nonzero even in a strictly Darwinian lineage, since acquired structure influences *who reproduces*. Mutual information cannot isolate direct write-back — directness is causal, not correlational — so WP0058 conditions on the selection signal, $\lambda_B := I_K(w : H' | H, \sigma)$ on the tagged write-back message $w = C(a)$, which attaches the quantity to the channel rather than inferring it from residual dependence, and blocks $a \rightarrow \sigma \rightarrow H'$ (`writeBack`). No additive split $\lambda_{\text{sel}} = I_K(a : H' | H) - \lambda_B$ is defined: conditional mutual information is not additively decomposable. Then: in the Darwinian regime there is no message, so $\lambda_B \leq O(\log)$ (`no_message_no_write_back`) — the premise is the *absence of the message*, architectural and posited, since no information measure certifies a missing causal path — while the *total* information crossing is bounded by the selection signal, $I_K(a : H' | H) \leq K(\sigma) + O(\log)$ (`darwinian_bandwidth_le_selection`, zeroth-order): that bound is the substantive half. In the Lamarckian regime $\lambda_B \leq K(w | H, \sigma) + O(\log)$ (`lamarckian_bandwidth_le_decoder_image`, first-order) — immediate, since λ_B is *defined* on w . `toyWB_writeback_zero_but_total_positive` machine-checks that the two quantities genuinely differ: a witness with $\lambda_B = 0$ and $\text{total} = 3$. Also `decoder_charged` ($K(C) \leq K(H) + O(\log)$, Prop. 1's AIT half) and `trivial_decoder_transmits_nothing` (Corollary 1: a trivial decoder transmits nothing however much was acquired). Witness `ToyWB` satisfies both named AIT hypotheses *and attains* the Darwinian bound (`toyWB_selection_bound_tight`) — the corollaries are tight, not merely non-vacuous.

- `persistence_conservation / conservation_tradeoff` — WP0162 Proposition 2: the ledger $K(O_W) = I_K(O_W : R) + K(O_W | R^*)$ (from symmetry of information), and at fixed $K(O_W)$, maximizing $I_K =$ minimizing the conditional residual.
- `meta_persistence / meta_persistence_limit` — WP0162 §6 (meta-persistence): a collective submodel of stable complexity $k > 0$ with a bounded transient $K(S_t^C | (S_{t+\tau}^C)^*) \leq L$ (delivered by closure + a Foster–Lyapunov contraction of the constituent flow) has collective persistence $\text{Pers}_C \geq 1 - (L + \text{slack})/k$, hence $\rightarrow 1$ as $k \rightarrow \infty$ — with *no* shared or compatible objective among the parts (the hypotheses mention only S^C, k, L, slack). Non-vacuity: `toy_meta_persistence_fires`.
- `regulator_selection` — WP0162 Proposition 1: $\arg \min_R K(R | W) = \arg \max_R [M(W : R) - K(R)]$.
- `regulator_selection_order` — the order form of the same proposition: under the chain rule, minimizing $K(R | W)$ is equivalent to maximizing $M(W : R) - K(R)$, as an order-equivalence rather than an argmin identity.
- `algorithmic_emergence` — WP0007, generic form: a correct general coarse-graining solver (one computing the minimal-sufficient-statistic structure function) is not computable. “Reduction is not construction.” The substance is Vereshchagin–Vitányi; `corollaryB` is the regulatory specialization.
- `AITProb.probabilistic_regulator_theorem_conditional` — ART sharp form in the $2^{-K(R|W)}$ reading (posterior favors regulators simple given the world).
- `PrefixMachine.lemma1_posterior_bounds` — the Bayes $\leftrightarrow$ Kolmogorov bridge, $\frac{1}{c_2} 2^{K(x)-|p|} \leq P(p | x) \leq \frac{1}{c_1} 2^{K(x)-|p|}$.
- `self_regulation_temporal_model` — Proposition 3, $K(S_{t+\tau} | S_t, E, C) \leq K(S_{t+\tau} | C) - \Delta_{\text{self}} + 2\text{slack}$.
- `pers_eq_nmai, persistent_pos` — persistence as temporal self-information.

Satisfiability witnesses: `toy_symmetry` (symmetry of information holds for Toy), `toyMachine_coding` (coding theorem holds for ToyMachine), `toy_self_model_fires` (the self-model corollary fires with a genuine gap $\Delta_{\text{self}} = 3$). **Guards:** `#check_failure` (peers A S) (part-whole), and `ystar_form_holds` vs `rawy_form_fails` (symmetry must use y^* , not y).

4 Theorem inventory and formalization status

Every numbered/named formal statement across the KT corpus, classified [STD] (standard AIT/probability/logic — axiom candidate) or [KT] (Kolmogorov-Theory-specific — proof target), with status. Sources: ART paper,² Pattern Persistence (WP0162),³ Agent Know Thyself (WP0192), Regulatory Coarse-Graining (WP0193), world-models paper,⁴ and the foundational AIT-of-consciousness paper.⁵

ID (paper)	Statement (abbreviated)	Class	Status
<i>ART paper — The Algorithmic Regulator (Entropy 2026, 28, 257)</i>			
Lemma 1	$P(p \mid x) = 2^{- p }/m(x)$; sandwiched by coding thm	STD	Proved (lemma1...)
Theorem 2	$\Pr((W, R) \mid x, E) \leq C 2^{M(W:R)} 2^{-\Delta}$	KT	Proved (prob_reg_thm)
Theorem 2 (sharp)	retains regulator cost: $\leq C 2^M 2^{-\Delta} 2^{-K(R)}$	KT	Proved (...sharp)
Lemma 2	OFF run lower-bounds world: $K(O_{W,\emptyset}) \leq K(W) + c_0$	KT	Axiom (hyp. hoff)
Eq. (6)	wrapper bound post $\leq (1/c_1) 2^{K(x)-K(W,R)}$	KT	Proved (wrapper_bound)
Theorem 1	coupled-pair posterior tilt $\in [\cdot] 2^{K(x)-K(W)-K(R)+M}$	KT	Proved (theorem1...)
Theorem 3	$m(O_{W,R})/m(O_{W,\emptyset}) \asymp 2^\Delta$ (as-if objective)	KT	Proved (theorem3...)
Theorem A4	low-complexity output $\Rightarrow$ strict tiny shrinkage of $K(W \mid R)$	KT	Proved (low_complexity_shrinkage)
Defs 1–2	internal model ($M(W:R) > 0$); good regulator ($\Delta > 0$)	KT	n/a (encoded)
Thms A1–A3, Lem A1	coding thm (uncond/cond), Kraft counting, excess-length decay	STD	Axiom
<i>Pattern Persistence — WP0162</i>			
Proposition 3	self-regulation $\Rightarrow$ temporal self-model (Eq. 44)	KT	Proved (self_reg...)
Proposition 1	regulator selection $\arg \min K(R \mid W) = \arg \max [M - K(R)]$	KT	Proved (regulator_selection)
Proposition 2	persistence conservation ledger $K(O_W) = I_K + K(O_W \mid R^*) \pm O(\log)$	KT	Proved (persistence...)
Meta-persistence	collective $\text{Pers}_C \geq 1 - (L+\text{slack})/k$; parts need not align	KT	Proved (meta_persistence)
Proposition 4	self-model incompleteness (3 obstructions)	KT	Proved (quine_floor,...)
Orbit-label thm	conserved orbit label = “generalized energy”	KT	Proved (genEnergy_conserved)
<i>Agent Know Thyself — WP0192</i>			
Proposition 1	temporal self-model (= WP0162 Prop. 3)	KT	Proved (same)

ID (paper)	Statement (abbreviated)	Class	Status
Principle 1	self-model incompleteness (= WP0162 Prop. 4)	KT	Proved (Phase 4)
Conjecture	acquired self-model prior	KT	n/a (conjecture)
<i>Regulatory Coarse-Graining — WP0193</i>			
Theorem B	targeted sufficient projection uncomputable	STD	Proved (theoremB)
Corollary B	regulatory coarse-graining uncomputable	KT	Proved (corollaryB)
Lemma 1	uncomputability of $h_{x \rightarrow y}$	STD	Axiom (via V-V)
Proposition 1	existence of a bounded regulation-sufficient projection	KT	Proved (existence)
Theorem A	selection uncomputability	KT	Proved (theoremA)
Cor. C1–C3	blanket uncomputability/fuzziness; acquisition; discovery	KT	n/a (interp.)
<i>Inheritance Architecture — WP0058</i>			
Proposition 1	no universal decoder (no total domain-recognizing inverse compiler)	KT	Proved (no_universal_decoder)
— engine	a total such inverse would decide achievability	KT	Proved (achievable_computable_of_inverse)
— non-vacuity	some computable D has an undecidable achievable set	STD	Axiom (ExistsUndecidableAchievableSet)
Prop. 1 (2nd half)	decoder charged to the program, $K(C_D) \leq K(H) + O(\log)$	KT	Proved (decoder_charged)
Master bound	$\lambda_B \leq K(H' H) + O(\log)$; the acquired state drops out	KT	Proved (bandwidth_le_cond)
Data processing	$I_K(a : z H) \leq K(z H) + O(\log)$; info via the channel only	KT	Proved (condIK_le_condRight)
Prop. 2 (Darwinian, λ_B)	write-back channel empty: $\lambda_B = I_K(a : H' H, \sigma) \leq O(\log)$	KT	Proved (no_message_no_write_back)
Prop. 2 (Darwinian, λ_{sel})	zeroth-order: total $I_K(a : H' H) \leq K(\sigma)$	KT	Proved (darwinian_bandwidth_le_selection)
Prop. 2 (Lamarckian)	first-order: $\lambda_B \leq K(C(a) H)$, the decoder image	KT	Proved (lamarckian_bandwidth_le_decoder_image)
— joint $\geq$ marginal	$K(x z) \leq K(\langle x, y \rangle z) + O(\log)$	STD	Axiom (JointGeMarginal)
— cond. subadditivity	$K(x) \leq K(x y) + K(y) + O(\log)$	STD	Axiom (SubadditivityCond)
— conditioning $\leq$ uncond.	$K(x y) \leq K(x) + O(\log)$	STD	Axiom (CondLeUncond)
Corollary 1	write-back cannot bootstrap novelty	KT	Proved (trivial_decoder_transmits_nothing)
Hypothesis 1	persistence matching, $\lambda_P^* \sim \tau_E$	KT	n/a (conjecture)
<i>World models / Lie groups — Entropy 2025 (separate geometry track)</i>			
Thms A1–A5	Lie transitivity, homogeneous space, moduli stack	STD	n/a (out of scope v1)
Noether (gen.)	continuous symmetry $\Rightarrow$ conserved quantity (abstract flow core)	KT	Proved (trajLabel_conserved)

ID (paper)	Statement (abbreviated)	Class	Status
Homogeneous space (transitive action $\Rightarrow G/H$)	—	STD	Proved (homogeneousSpace)
Lie pseudogroups / moduli stacks / Mostow / dim formula	—	geom	TODO (not in Mathlib)
<i>Foundational — nix019 / K-input</i>			
Hyp. 1–2, Cons. 1–2	simplicity selection; structured experience; high MAI	KT	n/a (informal)

Legend. **Proved**: machine-checked in Lean, no sorry. **Axiom**: assumed (standard result, or a named hypothesis consumed by a corollary). **TODO**: a KT proof target not yet formalized. n/a: definition, conjecture, interpretive remark, or informal claim (not a proof target in the current scope).

5 Roadmap

Priorities are ordered to maximize KT-correctness assurance per unit effort, exploiting the fact that the AIT/Bayes layer (§2) may be assumed.

Phase 1 — close the ART chain. DONE.

The wrapper bound Eq. (6) is now a *derived* theorem (wrapper_bound) from the coding theorem (CodingLB); ART **Theorem 1** (theorem1_posterior_tilt) and **Theorem 3** (theorem3_onoff_evidence, log-free form) are proved; probabilistic_regulator_theorem now consumes wrapper_bound and rests on the coding theorem + Lemma 2. All four #print axioms = Lean core only.

Phase 2 — persistence conservation (WP0162 Prop. 2). DONE.

Ledger $K(O_W) = I_K(O_W:R) + K(O_W \mid R^*)$ (persistence_conservation, from symmetry of information) and the trade-off “max $I_K = \min$ conditional residual” (conservation_tradeoff).

Phase 3 — regulator selection (WP0162 Prop. 1). DONE.

$\arg \min_R K(R \mid W) = \arg \max_R [M(W:R) - K(R)]$ proved (regulator_selection) from the chain rule; the ART sharp form’s $2^{-K(R|W)}$ reading is probabilistic_regulator_theorem_conditional.

Phase 4 — self-model incompleteness. DONE.

The three independent obstructions: the quine floor (quine_floor, $|Q_A| \geq K(A)$), the self-prediction dichotomy (self_prediction_dichotomy, an axiom-free diagonalization), and the Chaitin ceiling on minimality (chaitin_blocks_minimality).

Phase 5 — coarse-graining uncomputability (WP0193). DONE.

Theorem B and Corollary B (regulatory coarse-graining uncomputable) proved by an abstract reduction to Vereshchagin–Vítányi (theoremB, corollaryB).

All five roadmap phases are complete. The WP0162 / WP0192 / WP0193 KT corollaries and the full probabilistic ART layer are machine-checked, sorry-free, resting only on Lean core plus the named AIT axiom layer of §2.

Later — geometry track.

The Lie-group/Noether layer of the world-models paper is a distinct development (differential geometry in Mathlib), deferred until the AIT/KT core is complete.

Each phase keeps the invariant: KT results `sorry-free`; `#print axioms` = Lean core + named hypotheses; new hypotheses discharged by an extended toy model.

6 Conclusion

We have a green, machine-checked Level-1/1.5 core: the probabilistic ART (Theorem 2), the Bayes $\leftrightarrow$ Kolmogorov bridge (Lemma 1), the temporal self-model (Proposition 3), and persistence, all proved from an explicit, satisfiable AIT/Bayes axiom layer, with the KT ontology typed so the whole-vs-part error cannot compile. This document is the living status tracker; the roadmap above closes the remaining KT corollaries paper-by-paper while keeping the soundness invariants intact.

References

- [1] L. de Moura and S. Ullrich, *The Lean 4 Theorem Prover and Programming Language*, CADE 2021; The mathlib Community, *The Lean Mathematical Library*, CPP 2020.
- [2] G. Ruffini et al., *An Algorithmic-Information-Theoretic Regulator Theorem*, Entropy **28**(3):257, 2026.
- [3] G. Ruffini, *Pattern Persistence* (BCOM WP0162), 2026.
- [4] G. Ruffini et al., *World models, Lie groups, and symmetry in Kolmogorov Theory*, Entropy **27**(1):90, 2025.
- [5] G. Ruffini, *An algorithmic information theory of consciousness*, Neuroscience of Consciousness **2017**(1):nix019.
- [6] M. Li and P. Vitányi, *An Introduction to Kolmogorov Complexity and Its Applications*, 4th ed., Springer, 2019.