

BARCELONA COMPUTATIONAL FOUNDATION

Working Paper WP0195

Pillar: Kolmogorov Theory (KT)

A Lean 4 Formalization of Kolmogorov Theory

Axioms, Theorems, and Project Status

Machine-checking the KT corollaries against an explicit AIT interface

Giulio Ruffini with Claude Opus 4.8

Barcelona Computational Foundation · Neuroelectrics, Barcelona

`giulio.ruffini@bcom.one`

July 2026

Abstract

We are building **KTAIT**, a Lean 4 + Mathlib formalization whose purpose is narrow and load-bearing: to *machine-check that the Kolmogorov-Theory (KT) corollaries follow logically from an explicit algorithmic-information-theory (AIT) interface*, and that the KT ontology is typed so that structural errors (e.g. comparing a whole pattern to its own part) cannot silently compile. We do *not* re-prove classical AIT: the standard AIT and probability results (the invariance theorem, the coding theorem, symmetry of information, Kraft’s inequality, Bayes’ rule, the Solomonoff–Levin universal semimeasure, the recursion theorem, Rice’s and Chaitin’s theorems, and the Vereshchagin–Vitányi structure-function results) are taken as a clearly delimited *axiom layer*. Every *KT* corollary is then *proved* from that layer with no **sorry**, and a toy model witnesses that the axiom layer is jointly satisfiable (so the corollaries are not vacuously true). This working paper states the methodology, fixes the axiom layer, inventories every theorem/lemma across the KT corpus with its current formalization status, and sets a prioritized roadmap.

Keywords: Kolmogorov Theory, algorithmic information theory, Lean 4, Mathlib, formal verification, algorithmic regulator theorem, persistence, self-model.

Classification: MSC 68Q30, 03B35, 68V15 / ACM F.4.1

Contents

1	Introduction and methodology	4
1.1	The one rule	4
1.2	Why this is worth doing	4
1.3	Soundness: named hypotheses, not global axioms	4
1.4	Scope	4
1.5	Citing the formalization	4
2	The standard AIT / probability axiom layer	5
3	Current status of the Lean development	5
4	Theorem inventory and formalization status	14
5	Roadmap	18
6	Conclusion	19

1 Introduction and methodology

KT has matured to the point where the dominant risk is no longer in the classical results it invokes, but in whether the *new* KT corollaries are *correctly typed* and *logically entailed* by those results. This project addresses exactly that risk with a proof assistant.¹

1.1 The one rule

The one rule

Standard AIT (and probability) theorems are allowed to be *axioms*. KT corollaries must be *proved* from those axioms with no **sorry**.

If a KT corollary needs a **sorry**, that is a real gap in the paper. If a classical AIT fact is an axiom, that is honest and expected: we are checking the *dependency structure*, not the AIT foundations.

1.2 Why this is worth doing

The bug class we most want to exclude is *ontological*. For example, writing $I_K(A : \mathcal{S})$ where $\mathcal{S} \subset A$ is structurally vacuous (a whole versus its own part); the correct object is *temporal*, $I_K(\mathcal{S}_t : \mathcal{S}_{t+\tau})$ or $K(\mathcal{S}_{t+\tau} \mid \mathcal{S}_t, E, C) \ll K(\mathcal{S}_{t+\tau} \mid C)$. In KTAIT the KT roles (*substrate*, *pattern*, *self-code*, *readout*, *regulator*, *time*) are *distinct types*; the only bridge to a common carrier is an explicit projection, so the bad version *fails to type-check*. The whole-vs-part error becomes a typed obligation, not a silent mistake.

1.3 Soundness: named hypotheses, not global axioms

A naïve encoding axiom $F : \text{forall } (F : \text{AITFrame}), P F$ asserting a law for *every* frame is *inconsistent* here: one can hand-build a frame violating it (e.g. `slack=0` with mismatched K) and derive `False`, which would make all corollaries vacuously true. We therefore state each AIT fact as a *named hypothesis* (a `Prop` about a frame) and let each corollary consume it. Consequently `#print axioms` on every KT corollary shows only Lean’s core axioms (`propext`, `Classical.choice`, `Quot.sound`) — strictly stronger than “rests on named AIT axioms.” The toy model (§3) then exhibits a concrete frame satisfying the hypotheses, witnessing joint satisfiability.

1.4 Scope

The current effort is **Level 1** (a typed, axiomatic KT/AIT interface) plus **Level 1.5** (the faithful *probabilistic* ART, including its Bayesian posterior and the coding-theorem bridge). Out of scope for now: a computability-backed K from a concrete universal machine (Level 2), and the differential-geometric (Lie-group / Noether) layer of the world-models paper,⁴ which is a separate track.

1.5 Citing the formalization

Papers in the KT corpus cite this development with two macros. `\lean{...}` marks incidental Lean text: a module filename, a definition, one of Lean’s core axioms. `\ktait{decl}` is a machine-checked-claim citation: the prose at that point asserts a result and names the KTAIT declaration that proves it, so which claims are machine-checked is a lookup rather than an interpretation — for a reader, and for any tool indexing the corpus. Names inside `\ktait` must be declarations; the guard rejects filenames, paths, and core axioms there. Both macros are enforced: papers registered in `docs/citing-papers.txt` have every cited name resolved against

the current declarations by `scripts/check_sync.sh`, which the pre-push hook and CI run, so renaming or restating a theorem cannot silently break a citing paper’s claim. The canonical preamble definition is `\newcommand{\ktait}[1]{\texttt{\small #1}}`.

2 The standard AIT / probability axiom layer

These are the results we *assume*. They are classical and not in doubt; collecting them here fixes precisely what the KT proofs are allowed to lean on. (Per K’s directive, starting from these as axioms — and similarly for Bayes — is acceptable and intended.)

Axiom layer A — complexity

Invariance theorem. K is machine-independent up to an additive $O(1)$.

Coding theorem (unconditional). $c_1 2^{-K(x)} \leq m(x) \leq c_2 2^{-K(x)}$, i.e. $-\log_2 m(x) = K(x) \pm O(1)$.

Coding theorem (conditional). $m(x \mid y) \asymp 2^{-K(x \mid y)}$.

Kraft–McMillan. $\sum_p 2^{-|p|} \leq 1$, hence $\sum_x m(x) \leq 1$.

Symmetry of information. $I_K(x:y) = K(x) + K(y) - K(x, y) = K(x) - K(x \mid y^*) + O(\log)$; chain rule $K(x, y) = K(x) + K(y \mid x^*) + O(1)$.

Multiplicity compression. $N_{\leq L}(x) \geq 2^r \Rightarrow K(x) \leq L - r + O(\log L)$.

Axiom layer B — probability / Bayes

Solomonoff–Levin universal semimeasure. $m(x) = \sum_{U(p)=x} 2^{-|p|}$; conditional $m(R \mid W) \asymp 2^{-K(R \mid W)}$.

Bayes’ rule + deterministic likelihood. prior $P(p) = 2^{-|p|}$, likelihood $P(x \mid p) = 1\{U(p) = x\}$, evidence $P(x) = m(x)$.

Excess-length decay. $\Pr\{|p| \geq K(x) + k \mid x\} \leq 2C 2^{-k}$.

Axiom layer C — computability / logic

Kleene’s second recursion theorem (quines); **Rice’s theorem**; **Chaitin’s incompleteness** (T proves no $K(x) > n$ for $n > c_T$; K uncomputable); **Vereshchagin–Vítányi**: the Kolmogorov structure function $h_x(\alpha)$ is uncomputable and minimal sufficient statistics are not locatable.

WP0007 consumes three of these in distinct forms, and `AlgorithmicEmergence` keeps them apart because they are not interchangeable: `KUncomputable` (K is not computable), `KThresholdUndecidable` (the predicate $K(x) < b$ is not *decidable* — it is *semidecidable*, which is the whole content of the improvement barrier), and `KNotApproximable` (no computable f satisfies $K \leq f \leq K + c$, which does *not* follow formally from the first as stated). `chaitin_certification_ceiling` consumes the Chaitin bound in the “certified lower bounds are capped” form.

3 Current status of the Lean development

The repository (<https://github.com/giulioruffini/KTAIT>) builds green on Lean 4.31.0 + Mathlib v4.31.0. Modules:

Module	Content	Key results
Ontology.lean	six typed KT roles + part-whole guard	SelfCode carries isSub
Basic.lean	AITFrame, IK, NMAI, condStar	symmetry as named Prop
Probability.lean	prefix machine, Bayes posterior, coding thm	Lemma 1 (bridge)
ART.lean	AITProb (posterior), Theorem 2	probabilistic regulator thm
Persistence.lean	Pers, TemporalSelfInfo	persistence lemmas
SelfModel.lean	DeltaSelf, Proposition 3	temporal self-model
SelfModelLimits.lean	quine floor, diagonalization, Chaitin ceiling	WP0192 Principle 1
RegulatorSelection.lean	ChainRule	WP0162 Prop. 1 (regulator selection)
CoarseGraining.lean	RegSuff, structure function	WP0193 Thms A/B
AlgorithmicEmergence.lean	counting bound, raw-length and threshold barriers	WP0007 Thms 1–3, Props. 1–2, C
Contrast.lean	contrast fibre posterior	WP0162 Q3 (Eq. 26)
ShiftInvariance.lean	CheapTransform, PairSymmetric	WP0202 App. A (the null-hypothesis)
PushPull.lean	Plant, Admissible, InwardOnBoundary	WP0186 Thms 1–2, Cor. 1, gap Lemma
ModelOrPay.lean	MAICapped; multiplicity witnesses	WP0203 Props. 1–2
OrbitLabel.lean	conserved orbit labels	WP0162 App. C
NoetherFlow.lean	trajLabel, flow symmetries	generalized Noether (abstract core)
Decoder.lean	Achievable, Inverts, dovetail	WP0058 Prop. 1(i)–(iii), axiom-free
WriteBack.lean	bandwidth, DataProcessingCond	WP0058 Prop. 2–3, channel bound
Retention.lean	J, qStar, FirstCrossing	WP0058 Prop. 4 (optimal retention)
ToyModel.lean	Toy, ToyMachine, ToyGap	satisfiability witnesses
PatternPersist.lean	WP0162 ontology on AITFrame	four cells are agents; Prop. 3 inheritance
BadStatements.lean	guards bite	part-whole; y vs y^* ; WP0162 ontology

Proved KT corollaries (no sorry; #print axioms = Lean core only):

- `PP.token_is_agent`, `PP.class_is_agent`, `PP.collective_is_agent`, `PP.kind_is_agent` — the four cells of the WP0207 square (token a_i , constituent class $\mathcal{T}$, collective M , collective kind $\mathcal{L}(M)$) are all agents. Possible only because agenthood is a *predicate* on the single carrier Pattern, following WP0162’s “no agent over and above the pattern”.
- `PP.token_not_telehomeostatic`, `PP.class_is_telehomeostatic` — telehomeostasis is fixed by the objective alone (WP0162 §3), so a token whose objective proxies its class is a genuine agent that is not telehomeostatic. `PP.agentoptosis_gap` states the conjunction: that gap is the room agentoptosis occupies (WP0207). `PP.token_usually_telehomeostatic` records the default case: when the objective targets the token itself, it *is* telehomeostatic, so the divergence is the exception and not the rule.
- `PP.collective_not_telehomeostatic` — the same divergence one grain up: a collective instance whose objective proxies its lineage, with `PP.kind_is_telehomeostatic` closing the square. This is the snowflake-yeast sign anomaly as a typing fact rather than an exception.
- `PP.persists_iff_nmai` — the module’s persistence predicate unfolds definitionally to WP0162’s $\text{NMAI}(S_t, S_{t+\tau})$. `PP.Persists` is `Persistence.Persistent` on the pattern’s trajectory, so persistence is inherited from `Persistence.lean` rather than restated: rational-valued, normalized, and lag-indexed.
- `PP.agent_has_temporal_self_model` — because `PP.IsAgent` is a positive self-regulation gap Δ_{self} (`SelfModel.DeltaSelf`), Proposition 3 applies to any agent verbatim: its future self-code is cheap given its present organization. Nothing is re-proved; the same named hypotheses are consumed.
- `PP.externally_regulated_not_agent`, `PP.agent_need_not_be_telehomeostatic` — agenthood

is computed from Δ_{self} , never asserted, so a pattern with no gap is refused agenthood; and an agent whose objective targets another pattern is not telehomeostatic.

- `PP.room_not_agent`, `PP.decoupled_not_agent` — locus is *derived* from a directional ablation pair, never asserted, so the externally regulated thermostat room and the decoupled pattern are refused agenthood by computation.
- `AITProb.probabilistic_regulator_theorem` — ART Theorem 2, $\Pr((W, R) \mid x, E) \leq C 2^{M(W:R)} 2^{-\Delta}$.
- `AITProb.wrapper_bound` — ART Eq. (6), now *derived* from the coding theorem: $\text{post}(e, x) \leq (1/c_1) 2^{K(x)-K(e)}$.
- `AITProb.probabilistic_regulator_theorem_sharp` — ART Theorem 2 *sharp form* retaining the regulator cost: $\text{post} \leq C 2^{M(W:R)} 2^{-\Delta} 2^{-K(R)}$ (no extra hypothesis; the headline drops $2^{-K(R)} \leq 1$ via `probabilistic_regulator_theorem_of_sharp`).
- `AITProb.theorem1_posterior_tilt` — ART Theorem 1, two-sided posterior tilt.
- `AITProb.theorem3_onoff_evidence` — ART Theorem 3, ON/OFF evidence ratio $\asymp 2^\Delta$ (log-free form).
- `CoarseGraining.theoremB / corollaryB` — WP0193 Theorem B / Corollary B: regulatory coarse-graining is uncomputable, by reduction to Vereshchagin–Vitányi.
- `quine_floor`, `self_prediction_dichotomy`, `chaitin_blocks_minimality` — WP0192 Principle 1 / WP0162 Proposition 4: the three independent obstructions to a complete, exact, certifiable self-model (the middle one is an axiom-free diagonalization).
- `contrast_posterior_ranks_by_complexity` — WP0162 Q3 (Eq. 26): the contrast-fiber posterior ranks by joint simplicity, with no $2^{-\Delta}$ tilt.
- `transform_preserves_IK`, `shift_is_not_a_null`, `cheap_transform_not_null` — WP0202 Appendix A. If $y = T_i(x)$ is computable from x^* plus a short index, then $I_K(x : T_i x) \geq K(T_i x) - K(\text{code } i) - 3 \text{ slack}$: a *cheaply-computable transform preserves algorithmic mutual information*. The corollary names an error that actually shipped into a draft — a **circular shift is not a null** for whole-object MAI, since $K(\sigma_s(x) \mid x^*) \leq K(s) + O(1) = O(\log n)$, so the shifted copy retains nearly all of its algorithmic information with the original. Using it as a surrogate subtracts the signal, and a real screening effect was consequently withdrawn as “indistinguishable from noise”. `cheap_transform_not_null` states the contradiction directly: no ε below $K(\sigma_s x) - K(s) - 3 \text{ slack}$ can satisfy `IsNullFor`. (A time-aligned mutual information *rate* is a different object, for which a shift *is* a valid null; that object is not an AIT quantity and is not formalized here. Keeping the two apart is the whole point.) Satisfiability witness: `ShiftFrame`, in which the bound is tight, $I_K(x : \sigma_s x) = K(x)$.
- `PushPull.theorem1_necessity` — WP0186 Theorem 1. Against a free two-sided disturbance ($\underline{D} < 0 < \overline{D}$), holding a band $[x_-, x_+]$ requires pull authority at the upper edge at least $\overline{D}$ and push authority at the lower edge at least $-\underline{D}$. Its contrapositive, `single_signed_cannot_regulate`, is the paper’s headline necessity claim: a single-signed controller cannot hold the band. Theorem 2 (`theorem2_sufficiency`) shows the saturating edge rule attains those conditions, and `regulation_iff` assembles the two halves into the equivalence the paper states in pieces; `inward_iff` is the step that collapses the $\forall D$ condition to its two worst cases. Formalizing them showed these are *not* dynamical theorems: their proofs evaluate the field only at x_- and x_+ , so their content is arithmetic on two boundary inequalities. The dynamics enter only

through Nagumo’s theorem, which is *assumed* (see below) and is never used to derive the sign conditions.

- `PushPull.nonstrict_fails_attraction` and its repair `margin_gives_uniform_decrease` — a defect in WP0186’s Theorem 2, and the fix. The first is a *counterexample*: a plant meeting every hypothesis of the sufficiency theorem (band $[-1, 1]$, disturbance range $[-1, 1]$, unit authorities, saturating rule) whose closed-loop field vanishes at $x = 5$ under the admissible realization $D \equiv \overline{D}$, so the constant trajectory $x(t) \equiv 5$ never enters the band. Invariance survives, ultimate boundedness fails: the escape clause “strict inequality whenever $D < \overline{D}$ ” is vacuous because $D \equiv \overline{D}$ is itself admissible. The repair supplies a strict off-band margin $\varepsilon > 0$, giving $\dot{x} \leq -\varepsilon$ and hence entry in time $\leq (x_0 - x_+)/\varepsilon$. The asymmetry is real: the necessary condition is tight and non-strict, while sufficiency for attraction needs one ε of slack.
- `PushPull.sign_diverse_card_ge_two`, `PushPull.minimal_is_two` — WP0186’s cardinality claim, with the distinction Lean forced into the open: the theorems bound *signs*, not counts. A regulating flow set must contain at least one flow of each sign, whence $|S| \geq 2$; and two flows, one of each sign, achieve it, so the bound is tight. Two pushes do not regulate however many there are.
- `PushPull.restoring_not_single_signed`, `PushPull.restoring_decomposition` — WP0186 Corollary 1. For $\kappa > 0$ the leaky-integrator field $-\kappa(x - x^*)$ is strictly positive below x^* and strictly negative above it, so no single sign-constrained flow realizes it; and it decomposes explicitly into a non-negative push active below x^* and a non-negative pull active above. The single-term restoring law is a push–pull motif in collapsed notation, not a counterexample to one.
- `PushPull.lemma1_gap_identity`, with its existential restatement `lemma1_gap_is_S_existential` — WP0186’s complexity-gap lemma, against the opaque `AITFrame` carrier. From (A2)+(A3) (`OffSplit`: $K(x_{\text{OFF}}) = K(D_{\text{pred}}) + K(\eta)$ within slack) and (A4) (`OnResidual`: $K(x_{\text{ON}}) = K(\eta)$ within slack), the gap is $\Delta = S \pm 2\text{slack}$ with $S := K(D_{\text{pred}})$: the residual $K(\eta)$, possibly $\Theta(N)$, cancels. The regulator’s dynamics play no role — the assumptions, not the mechanism, carry the lemma. `no_structure_no_gap` is the degenerate case: if the predictable component is trivial ($S = 0$) the gap vanishes up to slack, so a regulator that caps the amplitude of unpredictable disturbance claims no model content. **A retraction is recorded with these theorems.** An earlier draft chained the identity into ART to conclude that a sustained Δ *forces* model content $M(W:R) \gtrsim \Delta$. That inference is not licensed and the paper no longer makes it; the type-level tell is that I_K does not occur in the statement at all, and no arithmetic on `OffSplit/OnResidual` will introduce it. WP0186’s Proposition 1 is now a selection claim, which belongs with `RegulatorSelection.lean`.
- `ModelOrPay.simple_regulator_cannot_model` — WP0203 Proposition 1, from the single AIT fact that module assumes, `MAICapped` ($I_K(x:y) \leq K(y) + \text{slack}$, a named hypothesis about a frame). If $K(R) + \text{slack} < \Delta$ then $I_K(W:R) < \Delta$ for *every* world: a regulator simpler than the gap it opens cannot carry Δ bits of model. Contrapositively (`model_implies_big`), a pair satisfying the modelling disjunct $\Delta \leq I_K(W:R)$ has $\Delta \leq K(R) + \text{slack}$ — “model it” implies “be big.” The clamp then needs no special pleading: by `clamp_MAI_bounded_by_own_complexity`, $K(R) = O(1)$ alone caps its mutual algorithmic information with every world.
- `ModelOrPay.perpair_does_not_lift` and `multiplicity_cancels_penalty` — WP0203 Proposition 2, the multiplicity objection, as *witnesses* rather than assertions. For every k there is a family of 2^k terms each of weight at most 2^{-k} whose total is exactly 1; and $2^k \cdot 2^{-k} = 1$.

A per-pair bound therefore places no useful bound on a sum over a fiber of matching size — the ART per-pair penalty is real and is exactly consumed by the multiplicity. These are counterexamples to the lifting step; they establish no forward tail bound and are not evidence for one.

- `ModelOrPay.invertible_bound_insufficient` — a second *insufficiency* witness. The invertible-coupling inequality $I_K(W:R) \geq K(R) - K(x_{\text{ON}}) - O(\log)$ bounds in $K(R) - K(x_{\text{ON}})$, not in $\Delta = K(x_{\text{OFF}}) - K(x_{\text{ON}})$: the values $K(R) = 1$, $K(x_{\text{ON}}) = 0$, $K(x_{\text{OFF}}) = 100$, $I_K = 1$ satisfy it while leaving I_K far below the gap. The two bounds coincide exactly when $K(x_{\text{OFF}}) \leq K(R)$ (`invertible_reaches_gap_iff`), which is the hypothesis nothing supplies and a clamp violates. What is machine-checked is that the invertible-coupling bound is *insufficient*; nothing here establishes the positive claim it was meant to support.
- `OrbitLabel.genEnergy_conserved` — WP0162 App. C: generalized energy as a conserved orbit label (axiom-free).
- `boundary_sufficient` — WP0162 App. E: the thin boundary as a compressed sufficient statistic.
- `low_complexity_shrinkage` — ART Theorem A4; `CoarseGraining.theoremA`, `CoarseGraining.existence` — WP0193 Theorem A and Proposition 1.
- `Noether.trajLabel_conserved`, `Noether.conserved_comp_symm` — the abstract core of generalized Noether (Structured Dynamics paper): the trajectory label is conserved by the flow, and symmetries commuting with the flow act on conserved quantities. (Continuous-time analogue of the orbit-label theorem.)
- `Homogeneous.homogeneousSpace` — a transitive symmetry action makes the state space a homogeneous space, $X \simeq G/H_x$ (Theorem A4’s algebraic core; the manifold/dimension layer remains future work).
- `Decoder.achievable_computable_of_inverse` / `Decoder.no_universal_decoder` — WP0058 Proposition 1(i)–(ii): a *total* computable write-back encoder that recognizes its own domain would decide membership in the achievable set of the developmental map; hence none exists (`exists_map_without_decoder`), and every realized write-back channel is a partial map whose validity domain must be supplied or learned rather than recognized by the encoder itself. Notable as the one KT corollary here that assumes *no* AIT fact at all — the engine (`achievable_computable_of_inverse`) is unconditional, and only the non-vacuity of the undecidability hypothesis enters as a named Prop (`ExistsUndecidableAchievableSet`). Satisfiability witness: `inverts_id`.
- `Decoder.partial_inverse_exists` / `Decoder.achievable_rePred` — the two positive halves of the same proposition, which locate exactly what the impossibility costs. Dovetailing search gives a *partial* inverse that returns a genuine preimage wherever it is defined, so inversion as such is available; what is lost is the ability to *refuse*, since the search diverges off the achievable set. Correspondingly the achievable set of a computable developmental map is recursively enumerable, which machine-checks the informal half of WP0058’s proof (“enumerate programs, run D , collect outputs”) and leaves only undecidability assumed.
- `WriteBack.channel_bound` — WP0058’s channel-bound lemma as a single named result covering both links, $I_K(a : H' \mid w) \leq I_K(a : z \mid w) \leq K(z \mid w) + O(\log)$. The right link is `condIK_le_condRight`; the left is data processing, assumed as `DataProcessingCond`. Cite this for the whole chain and `condIK_le_condRight` for the second link alone.

- `WriteBack.darwinian_of_dataProcessing` — a Darwinian regime is exactly the general data-processing hypothesis instantiated at the selection signal, so `Darwinian` assumes nothing beyond `DataProcessingCond` plus the caller's Markov premise that σ is the only route from a into H' .
- `WriteBack.bandwidth_le_cond`, the channel bound, and the regime corollaries — WP0058 Proposition 3. Writing λ_B for the write-back bandwidth (defined below on the write-back program, *not* as $I_K(a : H' | H)$), the master bound $\lambda_B \leq K(H' | H) + O(\log)$ holds but is weak in the Darwinian regime, where undirected mutation makes $K(H' | H)$ large without transmitting anything acquired. The sharp statement is data processing (`condIK_le_condRight`): since a reaches H' only through a channel variable z , $\lambda_B \leq I_K(a : z | H) \leq K(z | H) + O(\log)$. Crucially, the write-back bandwidth is *not* $I_K(a : H' | H)$: that quantity is nonzero even in a strictly Darwinian lineage, since acquired structure influences *who reproduces*. Mutual information cannot isolate direct write-back — directness is causal, not correlational — so WP0058 conditions on the selection signal, $\lambda_B := I_K(w : H' | H, \sigma)$ on the tagged write-back program $w = C(a)$, which attaches the quantity to the channel rather than inferring it from residual dependence, and blocks $a \rightarrow \sigma \rightarrow H'$ (`writeBack`). No additive split $\lambda_{\text{sel}} = I_K(a : H' | H) - \lambda_B$ is defined: conditional mutual information is not additively decomposable. Then: in the Darwinian regime there is no write-back program, so $\lambda_B \leq O(\log)$ (`no_program_no_write_back`) — the premise is the *absence of the write-back program*, architectural and posited, since no information measure certifies a missing causal path — while the *total* information crossing is bounded by the selection signal, $I_K(a : H' | H) \leq K(\sigma) + O(\log)$ (`darwinian_bandwidth_le_selection`, zeroth-order): that bound is the substantive half. In the Lamarckian regime $\lambda_B \leq K(w | H, \sigma) + O(\log)$ (`lamarckian_bandwidth_le_decoder_image`, first-order) — immediate, since λ_B is *defined* on w . `toyWB_writeback_zero_but_total_positive` machine-checks that the two quantities genuinely differ: a witness with $\lambda_B = 0$ and total = 3. Also `decoder_charged` ($K(C) \leq K(H) + O(\log)$, Prop. 2's background-free form) and `trivial_decoder_transmits_nothing` (Corollary 1: a trivial write-back program transmits nothing however much was acquired). Witness `ToyWB` satisfies both named AIT hypotheses *and attains* the Darwinian bound (`toyWB_selection_bound_tight`) — the corollaries are tight, not merely non-vacuous.
- `WriteBack.decoder_charged_rel` — WP0058 Proposition 2 proper, the background-relative form. `decoder_charged` charges the *whole* write-back apparatus to the heritable program, but part of it is supplied by a background B : the cell, the parent, the niche, a teacher, an external training pipeline. Under conditional subadditivity relative to a background (`SubadditivityCondRel`, a named Prop) an encoder reconstructible from H and B jointly (`RecoverableFromRel`) costs the lineage only $K(C | B) \leq K(H | B) + O(\log)$. The hypothesis is strictly weaker than `RecoverableFrom` (`toyWB_background_strictly_weaker`) and the bound is attained away from zero (`toyWB_decoder_charged_rel_tight`); `ToyWB` satisfies it (`toyWB_subadditivityCondRel`). What is *not* shown, here or anywhere in the module: that write-back cannot produce novelty, or that the encoder's representational axes were found by prior selection.
- `WriteBack.retention_le_program` — the retention curve $M(n) = I_K(w : H_{g+n} | H_g)$ (`retention`) is bounded at every lag by $K(w | H_g) + O(\log)$: what survives never exceeds what was sent. A corollary of `condIK_le_condRight`.
- `Retention.bangbang_optimal` — the optimal inherited-memory kernel is a *threshold rule*. With payoff $J = \sum_{n < N} q(n) v(n)$ linear in a box-constrained kernel $q : \mathbb{N} \rightarrow [0, 1]$ (J , `Admissible`) and $v(n)$ the net value of still acting on n -generation-old information, every

admissible kernel is dominated termwise by $q^*(n) = \mathbb{1}[v(n) > 0]$ (qStar, with qStar_of_pos, qStar_of_nonpos, qStar_admissible). If v is antitone the retained set is an initial segment (optimal_kernel_antitone_is_initial_segment), so the optimum is “retain for L generations, then drop”, and L is the *first crossing*: $q^*(n) = 1 \iff n < L$ (qStar_eq_one_iff_lt_first_crossing, qStar_eq_zero_of_first_crossing_le) for L satisfying FirstCrossing, which exists whenever v is somewhere nonpositive (exists_firstCrossing).

- **Retention.persistence_is_first_crossing** — the punchline. Instantiating $v(n) = R(n) - \kappa$ (netValue, antitone by netValue_antitone) with R the predictive information the environmental feature still carries at lag n and $\kappa > 0$ the per-generation cost of carrying it plus the mismatch loss, the optimum retains exactly the lags where $\kappa < R(n)$. Optimal persistence is therefore the first lag at which R crosses the cost threshold — WP0058’s $\tau_{\rho, \theta}$ with $\theta = \kappa$. The matching relation is the *solution* of a linear retention problem, not a modelling choice, and the threshold is set by the cost.
- **Retention.exponential_crossing** — for $R(n) = R_0 e^{-n/\tau}$ with $R_0, \tau > 0$ and $\kappa > 0$, $\kappa < R(n) \iff n < \tau \log(R_0/\kappa)$, so the optimal kernel is the indicator of that interval (qStar_exponential) and the crossing lag is *proportional* to the environmental correlation time at fixed cost ratio (exponential_persistence_proportional_to_tau). Non-vacuity: vTest_antitone, vTest_firstCrossing and toyRetention_qStar_beats_full exhibit an antitone net value whose optimum is a nontrivial initial segment and strictly beats full retention ($3 > 2$). What is *not* proved: that any real channel sits at this optimum, that the payoff is linear in q , that costs are constant in n , or that decay is exponential. The empirical matching claim remains Hypothesis 1.
- **persistence_conservation / conservation_tradeoff** — WP0162 Proposition 2: the ledger $K(O_W) = I_K(O_W:R) + K(O_W | R^*)$ (from symmetry of information), and at fixed $K(O_W)$, maximizing $I_K =$ minimizing the conditional residual.
- **meta_persistence / meta_persistence_limit** — WP0162 §6 (meta-persistence): a collective submodel of stable complexity $k > 0$ with a bounded transient $K(S_t^C | (S_{t+\tau}^C)^*) \leq L$ (delivered by closure + a Foster–Lyapunov contraction of the constituent flow) has collective persistence $\text{Pers}_C \geq 1 - (L + \text{slack})/k$, hence $\rightarrow 1$ as $k \rightarrow \infty$ — with *no* shared or compatible objective among the parts (the hypotheses mention only S^C, k, L, slack). Non-vacuity: toy_meta_persistence_fires.
- **regulator_selection** — WP0162 Proposition 1: $\arg \min_R K(R | W) = \arg \max_R [M(W:R) - K(R)]$.
- **regulator_selection_order** — the order form of the same proposition: under the chain rule, minimizing $K(R | W)$ is equivalent to maximizing $M(W:R) - K(R)$, as an order-equivalence rather than an argmin identity.
- **AlgorithmicEmergence.no_universal_compressor / no_universal_compressor_micro** — WP0007 Theorem 2 (improvement barrier) in the form that carries the intuition: no total computable compressor returns, for every string admitting a description shorter than itself, such a description. Asks nothing about optimality. Consumes CompressibleUndecidable, the undecidability of $\{x : K(x) < |x|\}$, which is a *different* classical fact from KThresholdUndecidable — a diagonal slice of an undecidable set need not be undecidable.
- **AlgorithmicEmergence.no_compression_improver_of_raw** — WP0007’s arbitrary-threshold corollary as the paper states it: derived from no_universal_compressor by specializing at $b = |x|$, and consuming *no* undecidability input of its own. no_compression_improver proves the same

conclusion directly from `KThresholdUndecidable`; that route is sound but its hypothesis is not independent.

- `AlgorithmicEmergence.threshold_undecidable_of_compressible` — why. Undecidability of the threshold relation $\{(x, b) : K(x) < b\}$ follows from undecidability of its diagonal slice $\{x : K(x) < |x|\}$ along the computable map $x \mapsto (x, |x|)$. The converse fails, which is why the raw-length statement needs its own proof; once the slice is settled the full relation costs nothing further. So the development rests on one classical input here, not two.
- `AlgorithmicEmergence.compressor_of_threshold` — the hierarchy, recorded as a lemma because it is easy to assert backwards. A threshold procedure specializes at $b = |x|$ to a universal compressor, so “no universal compressor” *implies* “no threshold procedure”. The raw-length statement is the stronger one and is **not** a corollary of the arbitrary- b theorem; the converse fails, since a procedure required to work only at $b = |x|$ is a weaker object than one required to work at every b .
- `AlgorithmicEmergence.no_compression_improver / no_compression_improver_micro` — WP0007 Theorem 2 (improvement barrier): no computable procedure returns, for every string (resp. micro-experiment) and threshold b , a description shorter than b whenever one exists. Strictly weaker a demand than optimality, and still impossible: validity already gives $\text{len}(A(x, b)) < b \Rightarrow K(x) < b$, so the threshold guarantee would make the undecidable predicate $K(x) < b$ decidable. The AIT input is `KThresholdUndecidable`. The companion observation — that the predicate is *semidecidable*, so open-ended search finds a shorter description whenever one exists and only fails to halt otherwise — is deliberately not formalized: it would need an operational semantics for programs, a halting-in- t -steps predicate and the dovetailing construction. The theorem rules out a *total* procedure that halts on every input while exploiting every available compression; semidecidability is the complementary positive fact about partial search, so nothing proved here rests on it.
- `AlgorithmicEmergence.identification_barrier` — WP0007 Theorem 3: no computable procedure returns, for every finite micro-experiment, a shortest program for its macrodata. “Reduction is not construction.” The reduction runs along the `Faithful` shift-and-read embedding of WP0007 Theorem 1; the AIT input is the uncomputability of K , assumed as the named hypothesis `KUncomputable`. `identification_barrier_conditional` is the same statement with complexity taken conditional on the microlaw and observation map; it is a variant kept for the development’s own sake, and is *not* what WP0007’s conditional-form corollary says — that is `conditional_complexity_uncomputable` above. WP0007’s numbering has shifted three times as results were inserted ahead of others, so entries here name results rather than number them.
- `AlgorithmicEmergence.no_additive_approximation` — WP0007 Proposition 1: relaxing “shortest” to “within a fixed additive constant c of shortest” does not restore computability, for any c . The AIT input here is `KNotApproximable`, which is *not* implied by `KUncomputable` as stated — it is the Berry argument run with slack, proved in WP0007 Proposition 1 and assumed here.
- `AlgorithmicEmergence.counting_bound / few_low_complexity_strings / most_states_incompressible` — WP0007 Theorem 1 (existence barrier), proved outright rather than assumed: distinct strings need distinct shortest programs, so fewer than 2^m strings have a shortest conditional program of length $< m$, and the complement is therefore more than a $1 - 2^{-d}$ fraction of the state space. The bound holds for every *fixed* conditioning string shared across the ensemble and does not depend on it — no choice of a common microlaw and observation map improves the fraction.

- `AlgorithmicEmergence.conditional_complexity_uncomputable` — WP0007’s conditional-form corollary: conditional complexity relative to (D_n, C_n, n) is not computable. The conditioning *varies with the input*, so this does not follow from uncomputability at a fixed conditioning string; the formalization carries the paper’s diagonal argument, with the counting witness and the constant-size search as its two named hypotheses.
- `AlgorithmicEmergence.chaitin_certification_ceiling` — WP0007 Proposition 2: certified complexity lower bounds are capped at a constant fixed by the theory and the reference machine, so no macromodel longer than that constant is certifiably shortest for any record. Same obstruction as `chaitin_blocks_minimality`, stated for macromodel optimality rather than self-models. Uncomputability is uniform over the family; unprovability is per-instance.
- `algorithmic_emergence` — **superseded**. This declaration stated WP0007’s former “no algorithm finds the optimal coarse-graining” result. WP0007 v0.3.0 withdrew that theorem: the objective is degenerate under a constant observation map, the target moves as C varies, and no reduction from the unrestricted structure function into the induced subclass was ever given. The declaration is retained only because `corollaryB` documents the regulatory case, where WP0193 fixes the target and the reduction does go through. Cite `identification_barrier` for WP0007.
- `AITProb.probabilistic_regulator_theorem_conditional` — ART sharp form in the $2^{-K(R|W)}$ reading (posterior favors regulators simple given the world).
- `PrefixMachine.lemma1_posterior_bounds` — the Bayes $\leftrightarrow$ Kolmogorov bridge, $\frac{1}{c_2} 2^{K(x)-|p|} \leq P(p \mid x) \leq \frac{1}{c_1} 2^{K(x)-|p|}$.
- `self_regulation_temporal_model` — Proposition 3, $K(S_{t+\tau} \mid S_t, E, C) \leq K(S_{t+\tau} \mid C) - \Delta_{\text{self}} + 2 \text{ slack}$.
- `pers_eq_nmai`, `persistent_pos` — persistence as temporal self-information.

Assumed bridges, and what is deliberately not formalized. `PushPull.lean` assumes one classical control-theory fact — *Nagumo’s theorem*, that for a locally Lipschitz field on a closed interval an inward-pointing boundary condition is equivalent to forward invariance — as the predicate `NagumoBridge`. This is the dynamical counterpart of not re-proving the coding theorem. It is an assumed bridge, not a KT result, and nothing in the module derives the sign conditions from it. `ModelOrPay.lean` assumes `MAICapped` and nothing else. The headline conjecture it was written to check — that a faithful regulator sustaining Δ_N either has $I_K(W:R) = \Omega(\Delta_N)$ or dissipates $\Omega(\Delta_N) k_B T \ln 2$ — is *not stateable* in `AITFrame`: Q_{diss} is a physical quantity and the frame has no bath, no temperature, and no notion of a physical realization of W or R . The same holds for the escape dichotomy, which quantifies over “any physical realization,” and for the saturation claim, which quantifies over sinks of finite capacity. We record the absence as a specification of the work remaining — it needs stochastic thermodynamics over a defined coupling — and not as a result. The known theorem in that shape is Shannon-form ($\langle W_{\text{ext}} \rangle \leq -\Delta F + k_B T I(W; R)$); whether it admits a single-sequence algorithmic strengthening with I_K in place of I is open.

Satisfiability witnesses: `toy_symmetry` (symmetry of information holds for Toy), `toyMachine_coding` (coding theorem holds for ToyMachine), `toy_self_model_fires` (the self-model corollary fires with a genuine gap $\Delta_{\text{self}} = 3$). **Guards:** `#check_failure` (peers A S) (part-whole), and `ystar_form_holds` vs `rawy_form_fails` (symmetry must use y^* , not y).

4 Theorem inventory and formalization status

Every numbered/named formal statement across the KT corpus, classified [STD] (standard AIT/probability/logic — axiom candidate) or [KT] (Kolmogorov-Theory-specific — proof target), with status. Sources: ART paper,² Pattern Persistence (WP0162),³ Agent Know Thyself (WP0192), Regulatory Coarse-Graining (WP0193), world-models paper,⁴ and the foundational AIT-of-consciousness paper.⁵

ID (paper)	Statement (abbreviated)	Class	Status
<i>ART paper — The Algorithmic Regulator (Entropy 2026, 28, 257)</i>			
Lemma 1	$P(p x) = 2^{- p }/m(x)$; sandwiched by coding thm	STD	Proved (lemma1...)
Theorem 2	$\Pr((W, R) x, E) \leq C 2^{M(W:R)} 2^{-\Delta}$	KT	Proved (prob_reg_thm)
Theorem 2 (sharp)	retains regulator cost: $\leq C 2^M 2^{-\Delta} 2^{-K(R)}$	KT	Proved (.._sharp)
Lemma 2	OFF run lower-bounds world: $K(O_{W,\emptyset}) \leq K(W) + c_0$	KT	Axiom (hyp. hoff)
Eq. (6)	wrapper bound post $\leq (1/c_1) 2^{K(x)-K(W,R)}$	KT	Proved (wrapper_bound)
Theorem 1	coupled-pair posterior tilt $\in [\cdot] 2^{K(x)-K(W)-K(R)+M}$	KT	Proved (theorem1...)
Theorem 3	$m(O_{W,R})/m(O_{W,\emptyset}) \asymp 2^\Delta$ (as-if objective)	KT	Proved (theorem3...)
Theorem A4	low-complexity output $\Rightarrow$ strict tiny shrinkage of $K(W R)$	KT	Proved (low_complexity_shrinkage)
Defs 1–2	internal model ($M(W:R) > 0$); good regulator ($\Delta > 0$)	KT	n/a (encoded)
Thms A1–A3, Lem A1	coding thm (uncond/cond), Kraft counting, excess-length decay	STD	Axiom
<i>Pattern Persistence — WP0162</i>			
Proposition 3	self-regulation $\Rightarrow$ temporal self-model (Eq. 44)	KT	Proved (self_reg...)
Proposition 1	regulator selection $\arg \min K(R W) = \arg \max [M - K(R)]$	KT	Proved (regulator_selection)
Proposition 2	persistence conservation ledger $K(O_W) = I_K + K(O_W R^*) \pm O(\log)$	KT	Proved (persistence...)
Meta-persistence	collective $\text{Pers}_C \geq 1 - (L + \text{slack})/k$; parts need not align	KT	Proved (meta_persistence)
Proposition 4	self-model incompleteness (3 obstructions)	KT	Proved (quine_floor,...)
WP0162 ontology	four cells (token, class, collective, collective kind) are all agents	KT	Proved (PP.token_is_agent,...)
WP0162 §3	telehomeostasis is fixed by the objective alone	KT	Proved (PP.class_is_telehomeostatic)
WP0207	agentoptosis gap: $\text{agent} \wedge \text{not telehomeostatic}$	KT	Proved (PP.agentoptosis_gap)
WP0162 §5	agenthood is a positive Δ_{self} , computed not asserted	KT	Proved (PP.externally_regulated_not_agent)
WP0162 §4	persistence unfolds to NMAI, inherited from Persistence	KT	Proved (PP.persists_iff_nmai)
Proposition 3 (PP)	every agent has a temporal self-model	KT	Proved (PP.agent_has_temporal_self_model)

ID (paper)	Statement (abbreviated)	Class	Status
WP0207	agent yet not telehomeostatic	KT	Proved (PP.agent_need_not_be_telehomeostatic)
Orbit-label thm	conserved orbit label = “generalized energy”	KT	Proved (genEnergy_conserved)
<i>Agent Know Thyself — WP0192</i>			
Proposition 1	temporal self-model (= WP0162 Prop. 3)	KT	Proved (same)
Principle 1	self-model incompleteness (= WP0162 Prop. 4)	KT	Proved (Phase 4)
Conjecture	acquired self-model prior	KT	n/a (conjecture)
<i>Regulatory Coarse-Graining — WP0193</i>			
Theorem B	targeted sufficient projection uncomputable	STD	Proved (theoremB)
Corollary B	regulatory coarse-graining uncomputable	KT	Proved (corollaryB)
Lemma 1	uncomputability of $h_{x \rightarrow y}$	STD	Axiom (via V-V)
Proposition 1	existence of a bounded regulation-sufficient projection	KT	Proved (existence)
Theorem A	selection uncomputability	KT	Proved (theoremA)
Cor. C1–C3	blanket uncomputability/fuzziness; acquisition; discovery	KT	n/a (interp.)
<i>Inheritance Architecture — WP0058</i>			
Prop. 1(i)	no total, domain-recognizing write-back encoder	KT	Proved (no_universal_decoder)
— engine	a total such inverse would decide achievability	KT	Proved (achievable_computable_of_inverse)
Prop. 1(ii)	some computable D admits no such encoder	KT	Proved (exists_map_without_decoder)
— non-vacuity	some computable D has an undecidable achievable set	STD	Axiom (ExistsUndecidableAchievableSet)
Prop. 1(iii)	a <i>partial</i> recursive inverse exists, with domain exactly $\mathcal{A}$	KT	Proved (partial_inverse_exists, dovetail)
— $\mathcal{A}$ is r.e.	the achievable set of a computable D is recursively enumerable	KT	Proved (achievable_rePred)
— unbounded in time	the dovetailing search has no bound on its running time	KT	n/a (resource claim; no cost model in this development)
Prop. 2	heritable burden net of background, $K(C B) \leq K(H B) + O(\log)$	KT	Proved (decoder_charged_rel)
— background-free form	$K(C_D) \leq K(H) + O(\log)$	KT	Proved (decoder_charged)
— cond. sub-add. rel. B	$K(x B) \leq K(x \langle y, B \rangle) + K(y B) + O(\log)$	STD	Axiom (SubadditivityCondRel)
Channel bound (lemma)	both links: $I_K(a : H' w) \leq I_K(a : z w) \leq K(z w) + O(\log)$	KT	Proved (channel_bound)
— left link	data processing under the Markov premise $a \rightarrow z \rightarrow H'$ given w	STD	Axiom (DataProcessingCond)

ID (paper)	Statement (abbreviated)	Class	Status
— right link	$I_K(a:z w) \leq K(z w) + O(\log)$; info via the channel only	KT	Proved (condIK_le_condRight)
Master bound	$\lambda_B \leq K(H' H) + O(\log)$; the acquired state drops out	KT	Proved (bandwidth_le_cond)
Prop. 3 (Darwinian, λ_B)	write-back channel empty: $\lambda_B = I_K(a: H' H, \sigma) \leq O(\log)$	KT	Proved (no_program_no_write_back)
Prop. 3 (Darwinian, λ_{sel})	zeroth-order: total $I_K(a: H' H) \leq K(\sigma)$	KT	Proved (darwinian_bandwidth_le_selection)
— regime from data proc.	Darwinian is DataProcessingCond at the selection signal	KT	Proved (darwinian_of_dataProcessing)
Prop. 3 (Lamarckian)	first-order: $\lambda_B \leq K(C(a) H)$, the decoder image	KT	Proved (lamarckian_bandwidth_le_decoder_image)
— joint marginal	$K(x z) \leq K(\langle x, y \rangle z) + O(\log)$	STD	Axiom (JointGeMarginal)
— cond. sub-additivity	$K(x) \leq K(x y) + K(y) + O(\log)$	STD	Axiom (SubadditivityCond)
— condition- ing $\leq$ uncond.	$K(x y) \leq K(x) + O(\log)$	STD	Axiom (CondLeUncond)
Corollary 1	write-back is capped by the write-back program the encoder can form	KT	Proved (trivial_decoder_transmits_nothing)
Retention bound	$M(n) = I_K(w: H_{g+n} H_g) \leq K(w H_g) + O(\log)$	KT	Proved (retention_le_program)
Prop. 4(i)	linear box-constrained kernel: the optimum is bang-bang	KT	Proved (bangbang_optimal)
Prop. 4(ii)	antitone net value $\Rightarrow$ retained set is $[0, L)$	KT	Proved (optimal_kernel_antitone_is_initial)
— L is first crossing	$q^*(n) = 1 \iff n < L$ for FirstCrossing	KT	Proved (qStar_eq_one_iff_lt_first_crossing)
— existence of L	a first crossing exists if v is somewhere non-positive	KT	Proved (exists_firstCrossing)
— threshold rule	q^* well defined and admissible	KT	Proved (qStar_of_pos, qStar_of_nonpos, qStar_admissible)
Prop. 4(iii)	$v = R - \kappa$: optimum retains exactly where $\kappa < R(n)$, i.e. $\tau_{\rho, \theta}$	KT	Proved (persistence_is_first_crossing, netValue_antitone)
Prop. 4(iv)	$\kappa < R_0 e^{-n/\tau} \iff n < \tau \log(R_0/\kappa)$	KT	Proved (exponential_crossing, qStar_exponential)
— proportionality	crossing lag $\propto \tau$ at fixed cost ratio	KT	Proved (exponential_persistence_proportionality)
— non-vacuity	an antitone v whose optimum is a nontrivial initial segment	KT	Proved (vTest_antitone, vTest_firstCrossing)
Hypothesis 1	persistence matching, $\lambda_P^* \sim \tau_E$, in real lineages	KT	n/a (conjecture; the <i>optimization</i> is proved above)
<i>Push–Pull Motifs — WP0186</i>			
Theorem 1	opposed flows necessary: $H(x_+) \geq \overline{D}$ and $G(x_-) \geq -\underline{D}$	KT	Proved (theorem1_necessity)

ID (paper)	Statement (abbreviated)	Class	Status
— contrapositive	a single-signed controller cannot hold the band	KT	Proved (single_signed_...)
— boundary step	the $\forall D$ condition collapses to two worst cases	KT	Proved (inward_iff)
Nagumo's theorem	inward on $\partial B \Leftrightarrow$ forward invariance	STD	Axiom (NagumoBridge)
Theorem 2	saturating edge rule renders the band invariant	KT	Proved (theorem2_sufficiency)
Thms 1+2	sharp form: invariance $\Leftrightarrow$ both edge conditions	KT	Proved (regulation_iff)
Remark 2	non-strict off-band bound does <i>not</i> give attraction (counterexample)	KT	Proved (nonstrict_fails_...)
— repair	strict margin ε : $\dot{x} \leq -\varepsilon$, finite-time entry	KT	Proved (margin_gives_...)
Minimality	sign-diverse $\Rightarrow S \geq 2$; two flows attain it	KT	Proved (sign_diverse_..., minimal_is_two)
Corollary 1	$-\kappa(x - x^*)$ is not single-signed; splits into push + pull	KT	Proved (restoring_not_..., restoring_decomp...)
App. Lemma 1	gap identity $\Delta = S \pm O(\log N)$ from (A2)–(A4)	KT	Proved (lemma1_gap_...)
— (A2)–(A4)	OFF splits as $K(D_{\text{pred}}) + K(\eta)$; ON tracks η	STD	Axiom (OffSplit, OnResidual)
Remark 7	$S = 0 \Rightarrow$ no gap (pure noise buys no model content)	KT	Proved (no_structure_no_gap)
Prop. 1 (with-drawn)	sustained Δ forces $M(W:R) \gtrsim \Delta$	KT	n/a (retracted; not entailed)
<i>The ART Tail Does Not Follow — WP0203</i>			
Proposition 1	$I_K(W:R) \leq K(R) + O(\log)$: a cheap regulator cannot model	KT	Proved (simple_regulator_..., model_implies_big)
— clamp	$K(R) = O(1) \Rightarrow I_K(W:R) = O(1)$ for <i>every</i> world	KT	Proved (clamp_MAI_bounded_...)
— MAI cap	$I_K(x:y) \leq K(y) + O(\log)$	STD	Axiom (MAICapped)
Proposition 2	per-pair 2^{-k} does not lift to a fiber of size 2^k (witness)	KT	Proved (perpair_does_not_lift, multiplicity_...)
Invertible coupling	$I_K \geq K(R) - K(x_{\text{ON}})$ holds with $I_K \ll \Delta$ (witness of insufficiency)	KT	Proved (invertible_bound_...)
— when it reaches Δ	exactly when $K(x_{\text{OFF}}) \leq K(R)$	KT	Proved (invertible_reaches_...)
“Model it or pay for it”	$I_K = \Omega(\Delta_N)$ or $Q_{\text{diss}} = \Omega(\Delta_N)k_B T \ln 2$	KT	n/a (conjecture; not stateable in AITFrame)
<i>World models / Lie groups — Entropy 2025 (separate geometry track)</i>			
Thms A1–A5	Lie transitivity, homogeneous space, moduli stack	STD	n/a (out of scope v1)
Noether (gen.)	continuous symmetry $\Rightarrow$ conserved quantity (abstract flow core)	KT	Proved (trajLabel_conserved)

ID (paper)	Statement (abbreviated)	Class	Status
Homogeneous space (transitive action $\Rightarrow G/H$)	—	STD	Proved (homogeneousSpace)
Lie pseudogroups / moduli stacks / Mostow / dim formula	—	geom	TODO (not in Mathlib)
<i>Foundational — nix019 / K-input</i>			
Hyp. 1–2, Cons. 1–2	simplicity selection; structured experience; high MAI	KT	n/a (informal)

Legend. **Proved**: machine-checked in Lean, no sorry. **Axiom**: assumed (standard result, or a named hypothesis consumed by a corollary). **TODO**: a KT proof target not yet formalized. n/a: definition, conjecture, interpretive remark, or informal claim (not a proof target in the current scope).

5 Roadmap

Priorities are ordered to maximize KT-correctness assurance per unit effort, exploiting the fact that the AIT/Bayes layer (§2) may be assumed.

Phase 1 — close the ART chain. DONE.

The wrapper bound Eq. (6) is now a *derived* theorem (wrapper_bound) from the coding theorem (CodingLB); ART **Theorem 1** (theorem1_posterior_tilt) and **Theorem 3** (theorem3_onoff_evidence, log-free form) are proved; probabilistic_regulator_theorem now consumes wrapper_bound and rests on the coding theorem + Lemma 2. All four #print axioms = Lean core only.

Phase 2 — persistence conservation (WP0162 Prop. 2). DONE.

Ledger $K(O_W) = I_K(O_W:R) + K(O_W \mid R^*)$ (persistence_conservation, from symmetry of information) and the trade-off “max $I_K = \min$ conditional residual” (conservation_tradeoff).

Phase 3 — regulator selection (WP0162 Prop. 1). DONE.

$\arg \min_R K(R \mid W) = \arg \max_R [M(W:R) - K(R)]$ proved (regulator_selection) from the chain rule; the ART sharp form’s $2^{-K(R|W)}$ reading is probabilistic_regulator_theorem_conditional.

Phase 4 — self-model incompleteness. DONE.

The three independent obstructions: the quine floor (quine_floor, $|Q_A| \geq K(A)$), the self-prediction dichotomy (self_prediction_dichotomy, an axiom-free diagonalization), and the Chaitin ceiling on minimality (chaitin_blocks_minimality).

Phase 5 — coarse-graining uncomputability (WP0193). DONE.

Theorem B and Corollary B (regulatory coarse-graining uncomputable) proved by an abstract reduction to Vereshchagin–Vítányi (theoremB, corollaryB).

All five roadmap phases are complete. The WP0162 / WP0192 / WP0193 KT corollaries and the full probabilistic ART layer are machine-checked, sorry-free, resting only on Lean core plus the named AIT axiom layer of §2.

Later — geometry track.

The Lie-group/Noether layer of the world-models paper is a distinct development (differential geometry in Mathlib), deferred until the AIT/KT core is complete.

Each phase keeps the invariant: KT results `sorry-free`; `#print axioms` = Lean core + named hypotheses; new hypotheses discharged by an extended toy model.

6 Conclusion

We have a green, machine-checked Level-1/1.5 core: the probabilistic ART (Theorem 2), the Bayes $\leftrightarrow$ Kolmogorov bridge (Lemma 1), the temporal self-model (Proposition 3), and persistence, all proved from an explicit, satisfiable AIT/Bayes axiom layer, with the KT ontology typed so the whole-vs-part error cannot compile. This document is the living status tracker; the roadmap above closes the remaining KT corollaries paper-by-paper while keeping the soundness invariants intact.

References

- [1] L. de Moura and S. Ullrich, *The Lean 4 Theorem Prover and Programming Language*, CADE 2021; The mathlib Community, *The Lean Mathematical Library*, CPP 2020.
- [2] G. Ruffini et al., *An Algorithmic-Information-Theoretic Regulator Theorem*, Entropy **28**(3):257, 2026.
- [3] G. Ruffini, *Pattern Persistence* (BCOM WP0162), 2026.
- [4] G. Ruffini et al., *World models, Lie groups, and symmetry in Kolmogorov Theory*, Entropy **27**(1):90, 2025.
- [5] G. Ruffini, *An algorithmic information theory of consciousness*, Neuroscience of Consciousness **2017**(1):nix019.
- [6] M. Li and P. Vitányi, *An Introduction to Kolmogorov Complexity and Its Applications*, 4th ed., Springer, 2019.